

# GridStat on DETER Proposal

Ryan C. Goodfellow  
Information Sciences Institute  
Marina Del Rey, CA  
rgoodfel@isi.edu

August 10, 2012

## 1 Background

### 1.1 The Power System

In the modern world, we have come to enjoy electric power as a commodity for many decades now. The system that provides us with electricity, the electric power grid, began in 1882 when Thomas Edison launched the Pearl Street Station Power Plant. This first power grid provided electricity for several city blocks in New York city. Fast forward to today, the power grids that we have are so vast and interconnected that they commonly transcend our most significant political borders and are considered as a single transcontinental system. There have obviously been many engineering feats accomplished to bring us from the Pearl Street grid to what we see today. Electricity is a tricky commodity in that it can not be stored in significant capacities. This means that supply must always exactly meet demand in an electric power system. Electrical engineers have been able to achieve this in electric power systems by constructing instrumentation and control systems. In the early days this was relatively simple. System engineers would decide on some voltage level they wanted to keep the system at. The ratio of the actual system voltage to this decided operational voltage is called the per unit voltage of the electrical system. If the per unit voltage in the system started to ascend away from 1, then there was too much supply (generation) of electricity in the system. Dually, if the per unit voltage in the system started to descend away from 1 then there was not enough supply to meet demand (load). This is all based on the fundamental fact that  $P = VI \equiv V = P/I$  where P is power, I is current and V is voltage. If we want voltage stability at some target per unit quantity, in the face of fluctuating current demanded by the users of the system then we must control power generation in response to demand. On this basis, systems engineers could keep track of per unit voltages at carefully selected locations within a system, send these quantities to controllers at generation stations and keep the supply demand problem in check. As such systems have become more complex over the years we have given them a special name, *supervisory control and data acquisition* (SCADA) systems.

SCADA systems are only a part of the engineering makes electric power systems possible. Another significant area of engineering within the power system is what are known as *protection automation control* (PAC) systems. These systems protect the physical power system and its users from more stochastic sources of instability such as lightning storms, falling trees and backhoes. When lightning strikes a power line, a tree falls through one, or a backhoe digs through one, the devices that comprise PAC systems can detect the resulting electrical effect and take corrective action to prevent damage of equipment or injury to users. The canonical example of a simple PAC system is a protective relay and a breaker. A protective relay instrumenting a transmission line can detect when a tree falls through a transmission line by the electrical effect it has. Say the result of the tree falling on the line is the line falling to the ground and energizing the earth. This is known as a line to ground fault and will cause current in the line to shoot up and voltage to dip down, a condition called overcurrent. The relay on the transmission line can detect this condition. In response it opens a breaker, disconnecting the faulted line from its power source. If this relay were not present in such a situation, either a transformer connected to the line could burn out in a violent fashion due to extreme currents or in the worse case the generator supplying the line could become overloaded and break down. The

cost of the PAC system is disproportionately low to the potential damage caused in its absence. Without PAC systems stable operation of the electric power system would be impossible.

Together, SCADA and PAC systems provide for most of the stability and reliability that is seen in today's electric power systems. The evolution known as the 'Smart Grid' happening today is the process by which pure electrical and electromechanical systems within the SCADA and PAC realm are being replaced by digital microprocessor based devices interconnected by computer networks. With this movement comes a huge potential for advancement in power systems engineering. Realization of this potential, however, is not without cost. In the context of the power system, that cost is complexity. The fact that electric power is a commodity in modern economies is the result of over a century of interplay between the theory and practice of engineering electric power systems. Integrating microprocessor based devices operating over computer networks into the systems responsible ensuring stability and protection *increases the complexity* of the system as a whole. This engineering space is coming to be known as cyber-physical systems engineering. In the limit, the physical portion of a cyber-physical system stops at physical instrumentation and actuation and the cyber portion does everything else. To flesh this out a bit in context, for the power system, the only physical components of PAC and SCADA systems will be the devices instrumenting conducting equipment and the devices actuating breakers, switchgear, governors and other mechanical devices. Everything else will be handled by the cyber portion of the system. The motivation for this is programability. Unlike an electromechanical device, computers are arbitrarily programmable and thus flexible. A prime example of this is the modern digital relay. In the past, an electromechanical relay could detect overcurrent and operate a breaker, but that was about it. Electromechanical relays could not coordinate with one another to protect a transmission line *system*, or provide detailed instrumentation information about electrical events that induced control actions like today's protective relays. This example extends to many other components in the power system. Distributed generation such as wind and solar energy, which can only probabilistically provide power is a rising trend in many power systems. The stochastic nature of these sources can adversely affect the stability of the system. As more distributed generation sources proliferate into our power system, the need for programable distributed control will only be amplified. Digital relays operating over computer networks making possible distributed coordination within PAC systems was the first step. The state-of-the-art in coordination and control algorithms realized in this realm today are quite localized and very rudimentary. An example of this would be two protective relays at either end of a transmission line agreeing on the fact that there is a significant voltage differential across the line and that protective action must be taken. Distributed control (for the PAC realm) does not extend far beyond this level of complexity, in most cases due to the fact that the communication systems that exist can not support the desired algorithm. This is where we can bring field of *distributed computing* to the power system.

The integration of distributed systems into critical sections of the electric power system is quite possibly the grandest challenge that has ever faced coordinated distributed computing, in theory and practice. The magnitude of the time and space requirements are unprecedented. The power system runs at a rate of 60 Hz and flows at the speed of light, any distributed system operating in this environment must run at a scalar multiple of the rate and keep up with the flow. The size in space of the power system spans entire continents with digital participants at least on the order of  $10^5$  today and accelerating toward tomorrow.

## 1.2 GridStat

The power system is a rate based synchronous system. The control systems that operate in local and wide areas of the power grid instrument this rate based system and maintain its stability by performing asynchronous discrete control actions. Control systems operating in the wide area require rate based data delivery at a defined scalar multiple of the power system operational rate of 60 Hz. Over a decade of research has gone into the GridStat data delivery system to provide wide-area rate-based data delivery with the strong quality of service (QoS) guarantees demanded by control systems. GridStat is implemented as a data-centric publish/subscribe middleware. When subscriptions are made, GridStat guarantees that a data path will be constructed between the publisher and the subscriber that ensures the data rate and latency requirements of the subscription. GridStat also ensures that data paths from publishers to subscribers are fault tolerant

through routing algorithms that create disjoint optimal paths between publishers and subscribers. By virtue of being data-centric, publishers are not aware of subscribers and subscribers are not aware of publishers. Each participant is simply aware of the data it is providing or consuming, e.g. the data itself decouples its producers and consumers. The GridStat data delivery system exists as a few logical entities; publishers, subscribers, the data plane and the management plane. Publishers and subscribers publish and subscribe data respectively, the data and management planes work cooperatively between publishers and subscribers to get publication data to subscribers at the rate they require it. The data plane is a graph of entities called forwarding engines that forward data from publishers to subscribers. The management plane consists of QoS-Brokers that facilitate subscription requests by constructing paths in the data plane between publishers and subscribers that can guarantee the requested data-rate. Consider the figure below.

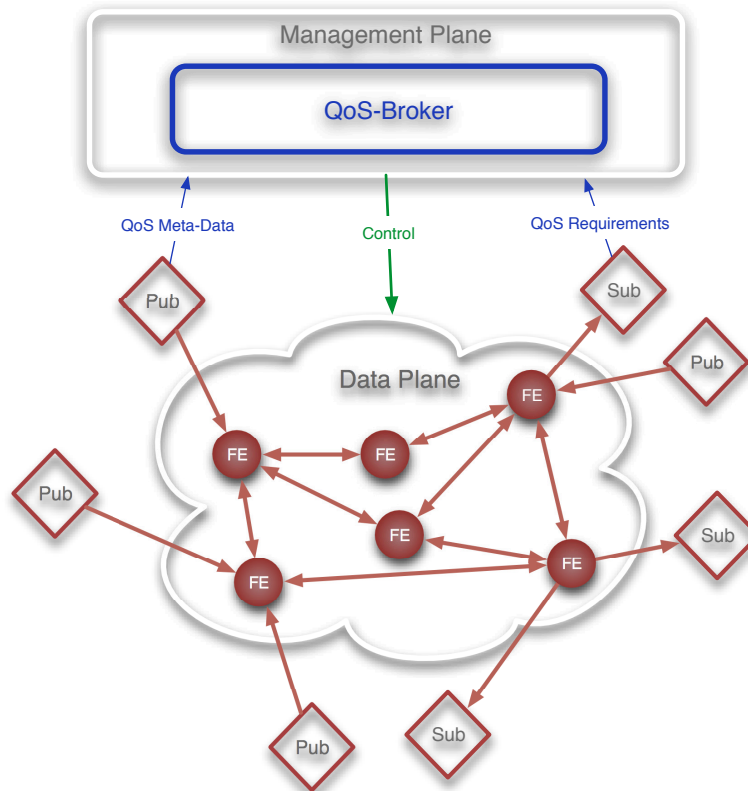


Figure 1: Simple GridStat System

Here we have a simple GridStat system with all of the entities just described. Let's walk through the basics of how the system works. When a publisher comes online it sends QoS meta-data to its QoS broker. This is basically making the QoS broker aware of both the presence of the data variable available for subscription and the maximum rate at which it may be provided. When a subscriber requests a variable at a given rate, it is the management plane's responsibility to construct a path through the data plane that gets the data from the publisher providing the variable to the subscriber requesting it, at the rate which the subscriber has requested it. If this is not possible then the management plane will reject the subscription request. The management plane will construct data paths in such a way that publishers and forwarding engines are only moving data at the minimum rate as dictated by subscriptions. For example, if a publisher is advertising a variable available at 1440 Hz, but the highest rate subscription for that request is at 30 Hz, than that variable will only move through the system at a maximum rate of 30 Hz. This concept also extends to path junctions in the data plane, if a forwarding engine is multiplexing a variable path downstream, it will only

provide each outgoing path with the minimum required data-rate. Thus, if a forwarding engine is receiving a variable stream at 60 Hz and is forwarding to 2 participants at 60 Hz and 25 Hz, then for the 25 Hz path, the incoming rate is down sampled to the required rate. This significantly improves the scalability of the system relative to common practices seen today in the power system such as 'wide-area ethernet' where all signals are sent at the highest rate possible. Notice that in this system there is only a single QoS-Broker. What this means is that this QoS-Broker 'owns' the entire data plane and all publications and subscriptions must go through this broker. Now consider the case where we have two QoS-Brokers.

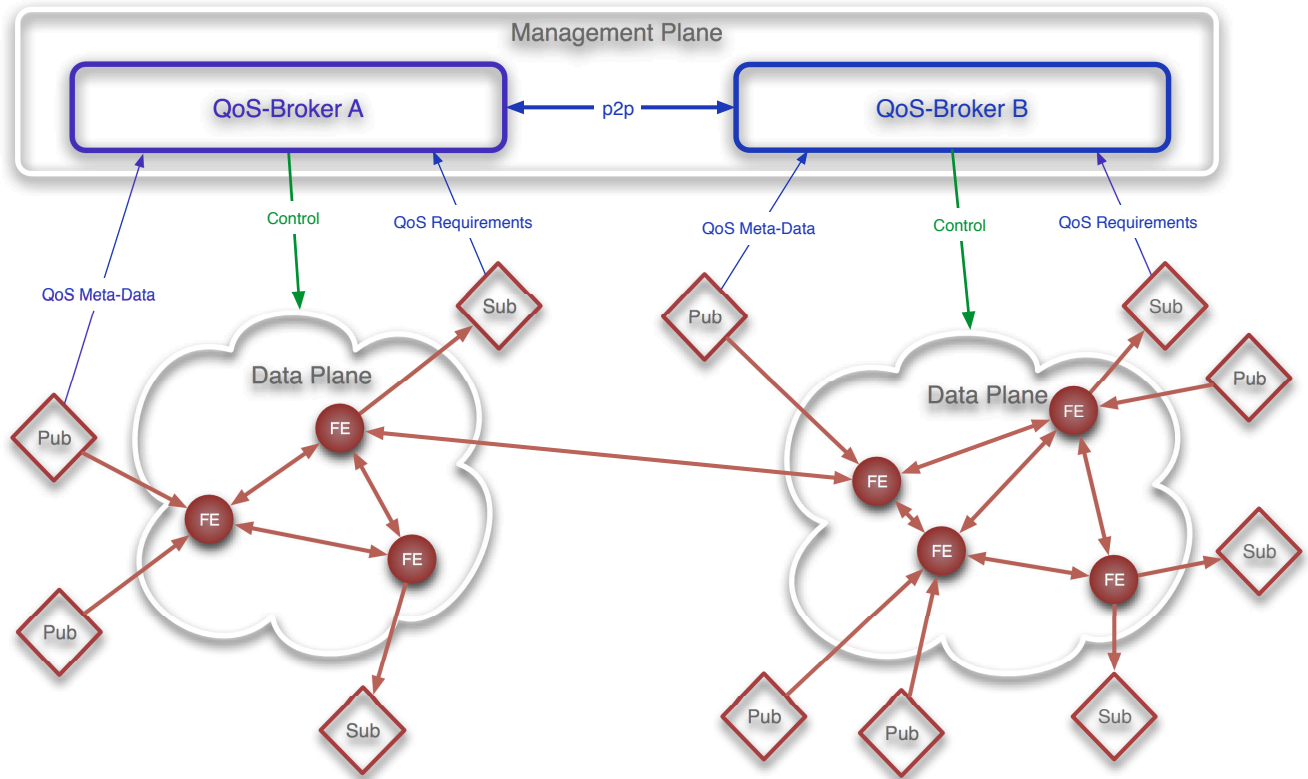


Figure 2: Two Broker GridStat System

The situation here is the same as the single broker except for the case when a subscriber in QoS-Broker B requests a subscription to data from a publisher in QoS-Broker A. In this case, each QoS-Broker must create the paths within its own data plane and they also must agree on a path between their respective data planes. There are many good reasons for partitioning the system into multiple brokers. In terms of performance and non-functional requirements; it can reduce the computational load experienced by the brokers from subscription requests, single points of failure for the entire system are at least reduced and individual logical brokers can be easily replicated for fault tolerance reasons. In terms of functional requirements, it allows for multiple participants in a GridStat system to cooperate in flexible and tightly managed ways. If the above diagram represented a transmission company (TRANSCO) and a distribution company (DISTCO), the TRANSCO certainly only wants to share certain information with the DISTCO and vice versa. Additionally each entity would want to control the number of external subscriptions and the rates at which those subscriptions are made available. Essentially the organizations are encapsulating the data within their network and protecting that network from being overloaded by external sources. In a peer-to-peer system of QoS-Brokers, this type of cooperation and control becomes possible.

Recent unfortunate history, with several major power-system black-outs shows us why wide-area situa-

tional awareness is an essential for maintaining system stability. This task requires robust real-time communication that transcends many corporate domains. GridStat aims to fill this role.

## 2 GridStat and DETER

In this section I will outline my proposal for deploying and experimenting with GridStat on the DETER testbed. The long-term goal of this work is to engineer a testbed environment within DETER for experimentation in wide-area communication and control for power systems. This problem space can be partitioned into three almost mutually exclusive areas of research and development; power system simulation, control system design and communication system design. Power system simulation exists on its own with no real dependency on either the communication or control system. The control system obviously must have hooks into the simulated power system to instrument signals and control mechanical actuation within the simulation, but that is about the extent of the interaction between the two. The complexity of power system simulators is very high, and unfortunately I do not know of any simulation systems that can take advantage of the scalability offered by cluster computing environments. Power system simulators today are designed for single processor systems, and consequentially, the answer to simulating a larger system is buying a faster processor (which is obviously a very limiting scalability factor). There are some simulators on the market that claim to support distributed computation, however, whether these systems have been proven with real large power systems is yet to be determined. Even if these simulators are ready for prime time, the immensely complex models that exist for real, large power systems must be translated into a format the distributed simulator can understand which itself is a substantial undertaking. In moving forward with this research, I strongly believe that monolithic simulator systems are *not* the answer. Thus research in this area will be necessary, however, because this area can be nicely decoupled from communication, which is my primary research area, it can wait.

The underlying requirements that motivated the original development of GridStat are from the power system, but are not exclusive to that domain. The essentials of the system are wide-area rate-based communication with real time requirements that span multiple domains of ownership. With this in mind, we can somewhat decouple GridStat from the power system and leverage DETER for experimentation and research of GridStat itself as a communication system that facilitates these goals. This is what I plan to focus on for the first phase of the research. However, keeping in mind the long term goals of this research, certain aspects will be influenced by the power system. This brings me to the deployment strategy for GridStat on DETER.

### 2.1 Deployment Topology

Several years ago, I lead a joint project between Washington State University (WSU), Bonneville Power Administration (BPA) and Schweitzer Engineering Laboratories (SEL). The goal of this project was building a GIS visualization system that facilitated the wide-area situational awareness made possible by synchrophasor data produced by SEL phasor measurement units (PMU). From this project I have gained a good working knowledge, as well as the actual data that describes the topology of the transmission system in the northwest region of the U.S. What this can provide for my research here, is a base GridStat system topology to start with that reflects a real world power system. Figure 3 shows a map of what this power transmission system looks like. What the map shows is blue dots representing substations and yellow lines representing transmission lines. In all there are 521 substations and 614 transmission lines owned by 117 different organizations. What this provides for us with the GridStat deployment on DETER is a nice reference model. We have 117 organizations, so we have at least 117 QoS-Brokers. There are far more organizations that own substations than own transmission lines. What this identifies is the set of substations at which the transmission system is feeding a distribution system or is being fed by a generation station.



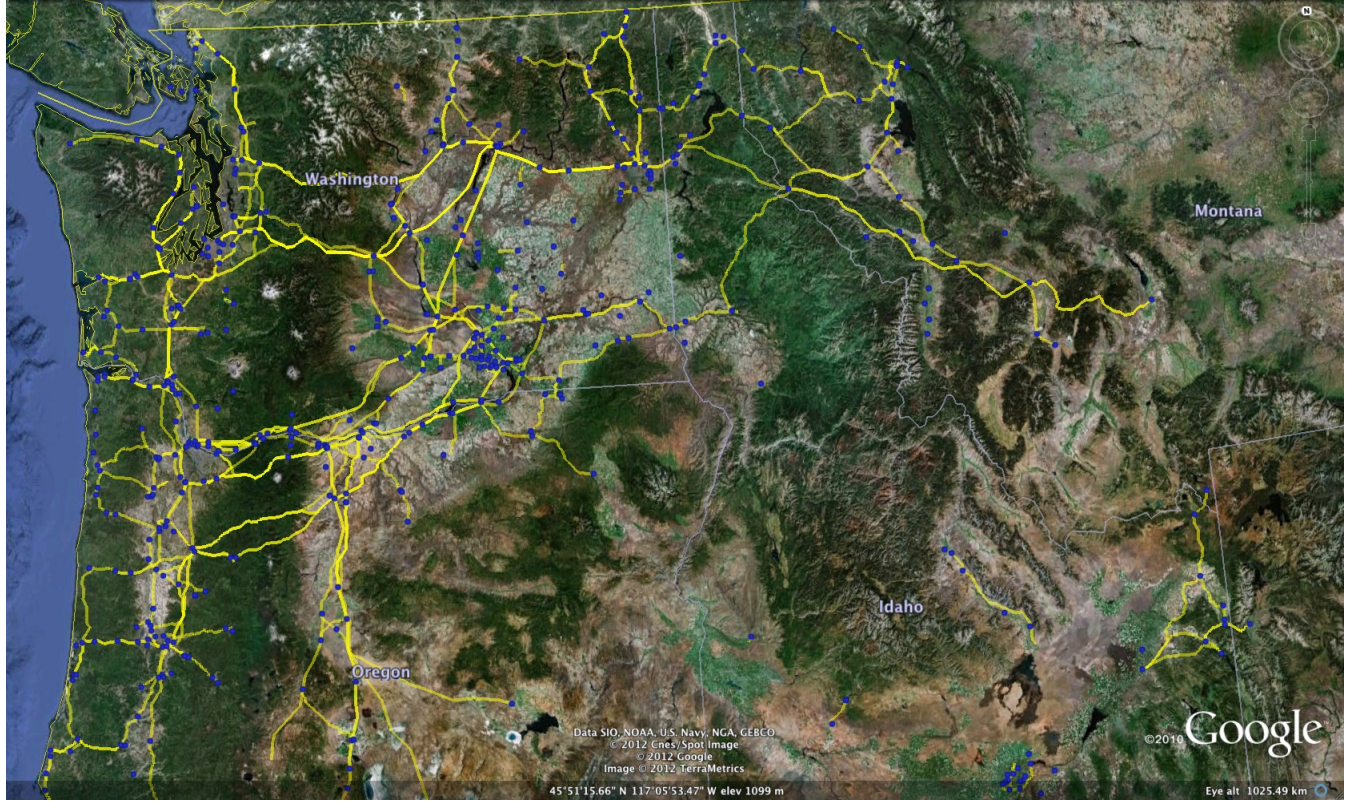


Figure 3: BPA Transmission System Topology

### 2.1.1 Publishers

For the initial deployment I propose that each transmission line will have a publisher at either end, providing mock synchrophasor data at a maximum rate of 60 Hz. This reflects the real world well from the perspective of the actual devices that are instrumenting the power lines and providing synchrophasor data. From an implementation perspective it is simple to produce the program that produces the mock synchrophasor data and provides it to GridStat. I actually think we have one of these in our code repos. For the substations that are owned by the distribution companies or generation companies, I propose that each have a single publisher that produces a synchrophasor that represents the low-voltage side of the substation. This basic infrastructure creates enough data to be interesting in the here-and-now for conducting experimentation on GridStat and for later down the road when real power system simulation and control comes into play. The GridStat infrastructure can be tested at the scale it is meant to operate at.

### 2.1.2 Subscribers

Transmission system subscribers will be broken down into two classes. First, wide-area control class subscribers will subscribe to all the data in the system at relatively low rates, say anything less than 10 Hz. While this may seem low it is much faster than rates realized by today's SCADA systems that operate of time intervals of a few seconds. The second class of subscribers for the TRANSCO will be regional subscribers. These subscribers will operate at rates closer to the publication maximum rate of 60 Hz, say between 30-60 Hz. The goal of these subscribers is to reflect the needs of regional controllers and state-estimators within the system. An example of this from the PAC realm could be a controller running an algorithm to protect a system of transmission lines running between a few substations. From the SCADA realm, there is active research currently happening on hierarchical state estimation at WSU that uses a similar approach with regional state estimators feeding a global state estimator to produce a global snap-

shot of the system. This approach not only reduces communications overhead, but also nicely distributes the computational overhead of the state estimation algorithm itself. The TRANSCO will also subscribe to DISTCO and GENCO publishers at a rate similar to the SCADA rate to keep track of load and generation characteristics. This will also help to flesh out some interesting details in federation policies between brokers.

From the distribution and generation side of things, each owning entity would subscribe to the publishers at each of its substations at variable rates. These entities would also subscribe to a small number of its nearest neighbor TRANSCO substations. I am not sure about the electrical motivations for this, however, from a communications standpoint, it provides a fault-tolerance mechanism for the DISTCO's and GENCO's with respect to the TRANSCO's. If for some reason a data stream fails from one TRANSCO forwarding engine to the DISTCO or GENCO it can resort to a nearest neighbor.

### 2.1.3 Data and Management Planes

As stated before the management plane will be partitioned as a QoS-Broker per owning entity, totaling in 117 brokers. This will be a scale never before tested for the management plane. Each QoS-Broker will be set up with a standard policy based on whether it is a TRANSCO broker or a DISTCO/GENCO broker. TRANSCO brokers and DISTCO/GENCO brokers will have to federate with one another in well defined ways. The policies set in place for each must allow the DISTCO/GENCO to acquire streams from the TRANSCO, while providing the TRANSCO with facilities to reject subscription requests from DISTCO/GENCO for information that it is not allowed to access. There is an analogous situation in the opposite direction from subscriptions propagating from the TRANSCO to the DISTCO/GENCO.

For the data plane, the publishers and subscribers are obviously co-located with their publishing and subscribing devices or applications. For forwarding engines we will select a set of substations along major transmission paths (500 kV and maybe 230 kV) to host forwarding engines. On the distribution and generation side, since we are not really representing these systems except for their connecting point with the transmission system they will just contain a single forwarding engine for each substation they own.

## 2.2 Implications

The deployment of the proposed GridStat implementation on DETER would allow for both scalability and security experimentation at a scale not seen before, specifically with the management plane. Tight management of the interactions with QoS-Brokers is very important for the potential users of GridStat. Security is of paramount concern here. A malicious broker or attacks on the management plane can have detrimental effects on the GridStat system. The DETER testbed provides an optimal environment to experiment with mechanisms engineered to defend against such circumstances at a scale that reflects the real target operating environment.

Although the 'electrical infrastructure' proposed is high level and simple, I think that it provides a good first top-down step toward total system representation. Applications including state estimation are possible at this high level. As research and development evolves in scalable power system simulation, this element can be incorporated into the framework along with the control systems that operate over it facilitated by the GridStat communications system.